

# Robust Optimisation Monte Carlo

Michael U. Gutmann

`michael.gutmann@ed.ac.uk`

Institute for Adaptive and Neural Computation  
School of Informatics, University of Edinburgh

22 November 2019

Borislav Ikonov and Michael U. Gutmann  
Robust Optimisation Monte Carlo  
*arXiv:1904.00670*, 2019

# Key messages

1. Optimisation Monte Carlo (OMC) is an existing method for efficient Bayesian inference with implicit models.
2. While efficient OMC under-estimates uncertainty by collapsing regions of near-constant likelihood into a single point.
3. A robust generalisation, robust OMC, explains and corrects this failure mode while maintaining OMC's benefits.

Background: Bayesian inference for implicit models

Contribution 1: A failure mode of Optimisation Monte Carlo

Contribution 2: Robust Optimisation Monte Carlo (ROMC)

Background: Bayesian inference for implicit models

Contribution 1: A failure mode of Optimisation Monte Carlo

Contribution 2: Robust Optimisation Monte Carlo (ROMC)

# Overall topic of the talk

Bayesian parameter inference for models where

1. the likelihood function is too costly to evaluate
2. exact sampling – simulating data from the model – is possible

(Diggle and Gratton, JRSS, 1982)

- ▶ Parametric statistical models specified by a data generating mechanism  $g : (\boldsymbol{\theta}, \mathbf{u}) \mapsto \mathbf{x} = g(\boldsymbol{\theta}, \mathbf{u})$ 
  - ▶  $\boldsymbol{\theta}$ : parameters
  - ▶  $\mathbf{u}$ : stochasticity / nuisance parameters with distribution  $p(\mathbf{u})$ .
  - ▶  $\mathbf{x}$ : generated data

(Diggle and Gratton, JRSS, 1982)

- ▶ Parametric statistical models specified by a data generating mechanism  $g : (\boldsymbol{\theta}, \mathbf{u}) \mapsto \mathbf{x} = g(\boldsymbol{\theta}, \mathbf{u})$ 
  - ▶  $\boldsymbol{\theta}$ : parameters
  - ▶  $\mathbf{u}$ : stochasticity / nuisance parameters with distribution  $p(\mathbf{u})$ .
  - ▶  $\mathbf{x}$ : generated data
- ▶ The (deterministic) function  $g$  and the distribution  $p(\mathbf{u})$  define the conditional distribution  $p(\mathbf{x}|\boldsymbol{\theta})$ 
  - ▶ evaluating  $p(\mathbf{x}|\boldsymbol{\theta})$  is generally intractable
    - likelihood function is intractable
  - ▶ drawing samples  $\mathbf{x}_i \sim p(\mathbf{x}|\boldsymbol{\theta})$  is possible
    - we can exploit this to perform Bayesian inference




- ▶ The function  $g(\boldsymbol{\theta}, \mathbf{u})$  is generally not known in closed form but implemented as computer code
  - ▶  $\boldsymbol{\theta}$ : input parameters
  - ▶  $\mathbf{u}$ : random draws performed when running the code / seed of the random number generator used

- ▶ The function  $g(\boldsymbol{\theta}, \mathbf{u})$  is generally not known in closed form but implemented as computer code
  - ▶  $\boldsymbol{\theta}$ : input parameters
  - ▶  $\mathbf{u}$ : random draws performed when running the code / seed of the random number generator used
- ▶ Other names: Simulator-based models, stochastic simulation model, generative (latent-variable) models, ...

# Importance

Such models and inference problems occur widely

- ▶ Evolutionary biology:  
Simulating evolution
- ▶ Neuroscience:  
Simulating neural activity
- ▶ Health science:  
Simulating the spread of an  
infectious disease
- ▶ Robotics:  
Simulating the outcome of an  
action
- ▶ Computer vision:  
Simulating naturalistic scenes
- ▶ ...



Simulated neural activity in rat somatosensory cortex  
(Figure from <https://bbp.epfl.ch/nmc-portal>)

# Bayesian inference for implicit models

- ▶ Task: Given

- ▶ observed data  $\mathbf{x}_o$ ,
- ▶ an implicit model  $g(\boldsymbol{\theta}, \mathbf{u})$ , and
- ▶ a prior distribution on  $\boldsymbol{\theta}$ ,

estimate the posterior  $p(\boldsymbol{\theta}|\mathbf{x}_o)$  / obtain approximate samples from it.

# Bayesian inference for implicit models

- ▶ Task: Given
  - ▶ observed data  $\mathbf{x}_o$ ,
  - ▶ an implicit model  $g(\boldsymbol{\theta}, \mathbf{u})$ , and
  - ▶ a prior distribution on  $\boldsymbol{\theta}$ ,

estimate the posterior  $p(\boldsymbol{\theta}|\mathbf{x}_o)$  / obtain approximate samples from it.

- ▶ Research fields: approximate Bayesian computation (ABC), likelihood-free inference, Bayesian indirect inference

(overviews: Sisson et al 2018, Lintusaari et al 2017, Gutmann and Corander 2016, Drovandi 2015, Marin et al 2012)

# Approximate Bayesian computation (ABC)

- ▶ ABC builds on the fact that samples from the posterior are given by samples from the prior for which simulated data  $\mathbf{x}$  are close to the observed data  $\mathbf{x}_o$ .

# Approximate Bayesian computation (ABC)

- ▶ ABC builds on the fact that samples from the posterior are given by samples from the prior for which simulated data  $\mathbf{x}$  are close to the observed data  $\mathbf{x}_o$ .
- ▶ Two core ingredients of ABC algorithms:
  1. a distance function  $d(\mathbf{x}, \mathbf{x}_o)$  between  $\mathbf{x}$  and  $\mathbf{x}_o$
  2. a search method to efficiently find such samples from the prior

▶ Ingredients:

- ▶ Distance  $d(\mathbf{x}, \mathbf{x}_o) = \|\Phi(\mathbf{x}) - \Phi(\mathbf{x}_o)\|_2$  with known  $\Phi$ .
- ▶ Search uses optimisation, which leads to increased efficiency.



- ▶ Ingredients:
  - ▶ Distance  $d(\mathbf{x}, \mathbf{x}_o) = \|\Phi(\mathbf{x}) - \Phi(\mathbf{x}_o)\|_2$  with known  $\Phi$ .
  - ▶ Search uses optimisation, which leads to increased efficiency.
- ▶ Assumptions:
  - ▶ (approximate) derivative of  $\Phi(\mathbf{x}) = \Phi(g(\boldsymbol{\theta}, \mathbf{u})) = \mathbf{f}(\boldsymbol{\theta}, \mathbf{u})$  wrt  $\boldsymbol{\theta}$  is available
  - ▶  $\dim(\boldsymbol{\theta}) \leq \dim(\Phi(\mathbf{x}))$

- ▶ Ingredients:
  - ▶ Distance  $d(\mathbf{x}, \mathbf{x}_o) = \|\Phi(\mathbf{x}) - \Phi(\mathbf{x}_o)\|_2$  with known  $\Phi$ .
  - ▶ Search uses optimisation, which leads to increased efficiency.
- ▶ Assumptions:
  - ▶ (approximate) derivative of  $\Phi(\mathbf{x}) = \Phi(g(\boldsymbol{\theta}, \mathbf{u})) = \mathbf{f}(\boldsymbol{\theta}, \mathbf{u})$  wrt  $\boldsymbol{\theta}$  is available
  - ▶  $\dim(\boldsymbol{\theta}) \leq \dim(\Phi(\mathbf{x}))$
- ▶ Algorithm to generate  $n$  weighted samples  $\boldsymbol{\theta}_i^*$ :
  - 1: **for**  $i \leftarrow 1$  to  $n$  **do**
  - 2:      $\mathbf{u}_i \sim p(\mathbf{u})$  ▷ Set seed
  - 3:      $\boldsymbol{\theta}_i^* = \arg \min_{\boldsymbol{\theta}} \|\mathbf{f}(\boldsymbol{\theta}, \mathbf{u}_i) - \Phi(\mathbf{x}_o)\|$  ▷ Optimisation
  - 4:     Compute  $\mathbf{J}_i$  with columns  $\partial \mathbf{f}(\boldsymbol{\theta}_i^*, \mathbf{u}_i) / \partial \theta_k$
  - 5:     Compute  $w_i = p(\boldsymbol{\theta}_i^*) * (\det(\mathbf{J}_i^\top \mathbf{J}_i))^{-1/2}$
  - 6:     Accept  $\boldsymbol{\theta}_i^*$  as posterior sample with weight  $w_i$

(Note that samples with too large final distances may be omitted.)

Background: Bayesian inference for implicit models

Contribution 1: A failure mode of Optimisation Monte Carlo

Contribution 2: Robust Optimisation Monte Carlo (ROMC)

# Application: Vision as inverse graphics

- ▶ Implicit model given by a graphics renderer
- ▶ We used Open Differential Renderer (Loper and Black, 2014)

20 parameters:

Shape

Rotation/Pose

Illumination

Colour


Renderer

(forward problem)



Inference

(inverse problem)



# Why Bayesian inference and not point-estimation?

- ▶ In some cases, quantifying uncertainty is very important
- ▶ The inverse problem may have multiple solutions (posterior may be multi-modal)

Example considered: Infer object colour when external lighting conditions are unknown.



(a) Gray teapot under red light.



(b) Red teapot under white light.

## Results for colour inference task

- ▶ We used OMC and the (simpler) rejection ABC algorithm.

## Results for colour inference task

- ▶ We used OMC and the (simpler) rejection ABC algorithm.
- ▶ Rejection ABC relies on trial and error instead of optimisation to determine the posterior samples. Slow but reliable.



## Results for colour inference task

- ▶ We used OMC and the (simpler) rejection ABC algorithm.
- ▶ Rejection ABC relies on trial and error instead of optimisation to determine the posterior samples. Slow but reliable.
- ▶ Same distance function  $d(\mathbf{x}, \mathbf{x}_o)$ : Euclidean distance between parameter predictions made by a neural network trained on images generated from the renderer under white light.



# Results for colour inference task

- ▶ We used OMC and the (simpler) rejection ABC algorithm.
- ▶ Rejection ABC relies on trial and error instead of optimisation to determine the posterior samples. Slow but reliable.
- ▶ Same distance function  $d(\mathbf{x}, \mathbf{x}_o)$ : Euclidean distance between parameter predictions made by a neural network trained on images generated from the renderer under white light.
- ▶ Posteriors for two colours  $c_0$  and  $c_1$  (red and green):




## Why did OMC fail?

- ▶ The OMC weights  $w_i = p(\theta_i^*) * (\det(\mathbf{J}_i^\top \mathbf{J}_i))^{-1/2}$  are unstable (ESS was 1.2!)
- ▶ This happens when the (approximate) likelihood function has nearly flat regions so that  $\det(\mathbf{J}_i^\top \mathbf{J}_i) \approx 0$ .

# Why did OMC fail?


- ▶ The OMC weights  $w_i = p(\theta_i^*) * (\det(\mathbf{J}_i^\top \mathbf{J}_i))^{-1/2}$  are unstable (ESS was 1.2!)
- ▶ This happens when the (approximate) likelihood function has nearly flat regions so that  $\det(\mathbf{J}_i^\top \mathbf{J}_i) \approx 0$ .



# Why did OMC fail?


- ▶ The OMC weights  $w_i = p(\theta_i^*) * (\det(\mathbf{J}_i^\top \mathbf{J}_i))^{-1/2}$  are unstable (ESS was 1.2!)
- ▶ This happens when the (approximate) likelihood function has nearly flat regions so that  $\det(\mathbf{J}_i^\top \mathbf{J}_i) \approx 0$ .

Note: stated OMC assumptions are not violated.




# Stabilising the weights/matrices does not help

- ▶ Taking the pseudo-inverse or pseudo-determinant of  $\mathbf{J}_i^\top \mathbf{J}_i$  does not help.



(a) Pseudo-inverse




(b) Pseudo-determinant


# Stabilising the weights/matrices does not help

- ▶ Taking the pseudo-inverse or pseudo-determinant of  $\mathbf{J}_i^\top \mathbf{J}_i$  does not help.
- ▶ The weights are not the real issue. The problem is more fundamental:

OMC uses a single point to represent an entire region where the likelihood is (nearly) constant.



(a) Pseudo-inverse



(b) Pseudo-determinant

Background: Bayesian inference for implicit models

Contribution 1: A failure mode of Optimisation Monte Carlo

Contribution 2: Robust Optimisation Monte Carlo (ROMC)

# Key properties of ROMC

(Ikonomov and Gutmann, *arXiv:1904.00670*, 2019)

1. **Fixes OMC's failure case:** It handles likelihood functions that are (nearly) flat on significant regions in parameter space.
2. Works for general distance functions  $d(g(\boldsymbol{\theta}, \mathbf{u}), \mathbf{x}_o)$  and not only Euclidean distances between summary statistics.  
(condition  $\dim(\boldsymbol{\theta}) \leq \dim(\Phi(\mathbf{x}))$  disappears)
3. Does not require (approximate) derivatives, while OMC does.
4. Can be run as post-processing to OMC or from scratch.



# The ROMC framework

ROMC is a framework for inference. It has three key steps:

1. For  $i = 1, \dots, n'$ , sample  $\mathbf{u}_i \sim p(\mathbf{u})$  and determine

$$\theta_i^* = \arg \min_{\theta} d(g(\theta, \mathbf{u}_i), \mathbf{x}_o)$$

Same as in OMC but we can use general distances  $d(\mathbf{x}, \mathbf{x}_o)$ .

# The ROMC framework

ROMC is a framework for inference. It has three key steps:

1. For  $i = 1, \dots, n'$ , sample  $\mathbf{u}_i \sim p(\mathbf{u})$  and determine

$$\theta_i^* = \arg \min_{\theta} d(g(\theta, \mathbf{u}_i), \mathbf{x}_o)$$

Same as in OMC but we can use general distances  $d(\mathbf{x}, \mathbf{x}_o)$ .

2. Use the minimal distances  $d_i^* = d(g(\theta_i^*, \mathbf{u}_i))$  to choose an acceptance threshold  $\epsilon$  / keep the  $n$  best  $\theta_i^*$ .

# The ROMC framework

ROMC is a framework for inference. It has three key steps:

1. For  $i = 1, \dots, n'$ , sample  $\mathbf{u}_i \sim p(\mathbf{u})$  and determine

$$\theta_i^* = \arg \min_{\theta} d(g(\theta, \mathbf{u}_i), \mathbf{x}_o)$$

Same as in OMC but we can use general distances  $d(\mathbf{x}, \mathbf{x}_o)$ .

2. Use the minimal distances  $d_i^* = d(g(\theta_i^*, \mathbf{u}_i))$  to choose an acceptance threshold  $\epsilon$  / keep the  $n$  best  $\theta_i^*$ .
3. For each  $i$  where  $d_i^* \leq \epsilon$ , define a proposal distribution  $q_i$  on the “acceptance region”  $C_\epsilon^i = \{\theta : d(g(\theta, \mathbf{u}_i), \mathbf{x}_o) \leq \epsilon\}$

# The ROMC framework

ROMC is a framework for inference. It has three key steps:

1. For  $i = 1, \dots, n'$ , sample  $\mathbf{u}_i \sim p(\mathbf{u})$  and determine

$$\theta_i^* = \arg \min_{\theta} d(g(\theta, \mathbf{u}_i), \mathbf{x}_o)$$

Same as in OMC but we can use general distances  $d(\mathbf{x}, \mathbf{x}_o)$ .

2. Use the minimal distances  $d_i^* = d(g(\theta_i^*, \mathbf{u}_i))$  to choose an acceptance threshold  $\epsilon$  / keep the  $n$  best  $\theta_i^*$ .
3. For each  $i$  where  $d_i^* \leq \epsilon$ , define a proposal distribution  $q_i$  on the “acceptance region”  $C_\epsilon^i = \{\theta : d(g(\theta, \mathbf{u}_i), \mathbf{x}_o) \leq \epsilon\}$

Approximate posterior is represented by weighted samples  $\theta_{ij}$ :

$$\theta_{ij} \sim q_i(\theta) \quad w_{ij} = \mathbb{1}_{C_\epsilon^i}(\theta_{ij}) \frac{p(\theta_{ij})}{q_i(\theta_{ij})} \quad (i=1, \dots, n; \quad j=1, \dots, m)$$

# Construction of the proposal distribution

(General idea, see paper for details)

- ▶ Using  $\theta_i^*$  and the optimisation trajectory, we build a model of the acceptance regions  $C_\epsilon^i = \{\theta : d(g(\theta, \mathbf{u}_i), \mathbf{x}_o) \leq \epsilon\}$

# Construction of the proposal distribution

(General idea, see paper for details)

- ▶ Using  $\theta_i^*$  and the optimisation trajectory, we build a model of the acceptance regions  $C_\epsilon^i = \{\theta : d(g(\theta, \mathbf{u}_i), \mathbf{x}_o) \leq \epsilon\}$
- ▶ Simple but effective: model  $C_\epsilon^i$  as a hypercube or ellipse and define  $q_i$  to be the uniform distribution on it.

# Construction of the proposal distribution

(General idea, see paper for details)

- ▶ Using  $\theta_i^*$  and the optimisation trajectory, we build a model of the acceptance regions  $C_\epsilon^i = \{\theta : d(g(\theta, \mathbf{u}_i), \mathbf{x}_o) \leq \epsilon\}$
- ▶ Simple but effective: model  $C_\epsilon^i$  as a hypercube or ellipse and define  $q_i$  to be the uniform distribution on it.
- ▶ Note: When computing the weight,

$$w_{ij} = \mathbb{1}_{C_\epsilon^i}(\theta_{ij}) \frac{p(\theta_{ij})}{q_i(\theta_{ij})}$$

the **indicator function** checks whether  $\theta_{ij}$  is in the true acceptance region  $C_\epsilon^i$ .  $\Rightarrow$  Some robustness to modelling errors.

# Construction of the proposal distribution

(General idea, see paper for details)

- ▶ Using  $\theta_i^*$  and the optimisation trajectory, we build a model of the acceptance regions  $C_\epsilon^i = \{\theta : d(g(\theta, \mathbf{u}_i), \mathbf{x}_o) \leq \epsilon\}$
- ▶ Simple but effective: model  $C_\epsilon^i$  as a hypercube or ellipse and define  $q_i$  to be the uniform distribution on it.
- ▶ Note: When computing the weight,

$$w_{ij} = \mathbb{1}_{C_\epsilon^i}(\theta_{ij}) \frac{p(\theta_{ij})}{q_i(\theta_{ij})}$$


the **indicator function** checks whether  $\theta_{ij}$  is in the true acceptance region  $C_\epsilon^i$ .  $\Rightarrow$  Some robustness to modelling errors.

- ▶ Check requires evaluating the distance  $d(g(\theta, \mathbf{u}_i), \mathbf{x}_o)$  and can be omitted/approximated to accelerate the inference.





# Results on the toy example

- ▶ Acceptance regions  $C_\epsilon^i$  given by intervals on the line.
- ▶ ROMC handles the (nearly) flat likelihood function correctly.
- ▶ ROMC accurately represents uncertainty while OMC does not.



# Results on the colour inference task

- ▶ Setup:
  - ▶ Optimisation via a gradient-free method (Bayesian optimisation with GP surrogate modelling)
  - ▶ Acceptance regions were modelled as ellipses (derived from the GP surrogate model)
- ▶ ROMC posterior matches reference posterior well.
- ▶ Effective sample size: 97% (vs. approx. 0.5% for OMC)



# ROMC generalises OMC (see paper for the proof)

Theorem: Under the below assumptions, ROMC becomes equivalent to standard OMC as  $\epsilon \rightarrow 0$ .

**Assumption 1.** The distance  $d(g(\boldsymbol{\theta}, \mathbf{u}), \mathbf{x}_o)$  is given by the Euclidean distance between summary statistics  $\|\mathbf{f}(\boldsymbol{\theta}, \mathbf{u}) - \Phi(\mathbf{x}_o)\|$ .

**Assumption 2.** The proposal distribution  $q_i(\boldsymbol{\theta})$  is the uniform distribution on  $C_\epsilon^i$ .

**Assumption 3.** The acceptance regions  $C_\epsilon^i$  are approximated by the ellipsoid  $C_\epsilon^i = \{\boldsymbol{\theta} : (\boldsymbol{\theta} - \boldsymbol{\theta}_i^*)^\top \mathbf{J}_i^\top \mathbf{J}_i (\boldsymbol{\theta} - \boldsymbol{\theta}_i^*) \leq \epsilon\}$  where  $\mathbf{J}_i$  is the Jacobian matrix with columns  $\partial \mathbf{f}(\boldsymbol{\theta}_i^*, \mathbf{u}_i) / \partial \theta_k$ .

**Assumption 4.** The matrix square root  $\mathbf{A}_i$  of  $\mathbf{J}_i^\top \mathbf{J}_i$  is full rank, i.e.  $\text{rank}(\mathbf{A}_i) = \dim(\boldsymbol{\theta})$ .

**Assumption 5.** The prior  $p(\boldsymbol{\theta})$  is constant on the acceptance regions  $C_\epsilon^i$ .

# Explanation of the failure case

Identified failure case is due to violation of Assumptions 3 and 4:

**Assumption 3.** The acceptance regions  $C_\epsilon^i$  are approximated by the ellipsoid  $C_\epsilon^i = \{\boldsymbol{\theta} : (\boldsymbol{\theta} - \boldsymbol{\theta}_i^*)^\top \mathbf{J}_i^\top \mathbf{J}_i (\boldsymbol{\theta} - \boldsymbol{\theta}_i^*) \leq \epsilon\}$  where  $\mathbf{J}_i$  is the Jacobian matrix with columns  $\partial \mathbf{f}(\boldsymbol{\theta}_i^*, \mathbf{u}_i) / \partial \theta_k$ .

**Assumption 4.** The matrix square root  $\mathbf{A}_i$  of  $\mathbf{J}_i^\top \mathbf{J}_i$  is full rank, i.e.  $\text{rank}(\mathbf{A}_i) = \dim(\boldsymbol{\theta})$ .

# Explanation of the failure case

Identified failure case is due to violation of Assumptions 3 and 4:

**Assumption 3.** The acceptance regions  $C_\epsilon^i$  are approximated by the ellipsoid  $C_\epsilon^i = \{\boldsymbol{\theta} : (\boldsymbol{\theta} - \boldsymbol{\theta}_i^*)^\top \mathbf{J}_i^\top \mathbf{J}_i (\boldsymbol{\theta} - \boldsymbol{\theta}_i^*) \leq \epsilon\}$  where  $\mathbf{J}_i$  is the Jacobian matrix with columns  $\partial \mathbf{f}(\boldsymbol{\theta}_i^*, \mathbf{u}_i) / \partial \theta_k$ .

**Assumption 4.** The matrix square root  $\mathbf{A}_i$  of  $\mathbf{J}_i^\top \mathbf{J}_i$  is full rank, i.e.  $\text{rank}(\mathbf{A}_i) = \text{dim}(\boldsymbol{\theta})$ .

For non-uniform priors, one then also risks violating Assumption 5:

**Assumption 5.** The prior  $p(\boldsymbol{\theta})$  is constant on the acceptance regions  $C_\epsilon^i$ .

- ▶ Talk was on Bayesian inference for implicit models.

- ▶ Talk was on Bayesian inference for implicit models.
  - ▶ Implicit models: models that are defined by a data generating process.

- ▶ Talk was on Bayesian inference for implicit models.
  - ▶ Implicit models: models that are defined by a data generating process.
  - ▶ Optimisation Monte Carlo (OMC): Bayesian inference method that uses optimisation to increase computational efficiency.



# Conclusions

- ▶ Talk was on Bayesian inference for implicit models.
  - ▶ Implicit models: models that are defined by a data generating process.
  - ▶ Optimisation Monte Carlo (OMC): Bayesian inference method that uses optimisation to increase computational efficiency.
- ▶ We showed that OMC under-estimates posterior uncertainty by collapsing regions of near-constant likelihood into a point.

# Conclusions

- ▶ Talk was on Bayesian inference for implicit models.
  - ▶ Implicit models: models that are defined by a data generating process.
  - ▶ Optimisation Monte Carlo (OMC): Bayesian inference method that uses optimisation to increase computational efficiency.
- ▶ We showed that OMC under-estimates posterior uncertainty by collapsing regions of near-constant likelihood into a point.
- ▶ We proposed a robust generalisation of OMC, robust OMC, that explains and corrects this failure mode while maintaining OMC's benefits due to optimisation.